
Advanced Alignment Models

Philipp Koehn

11 February 2016



IBM Model 1



- Generative model: break up translation process into smaller steps
 - IBM Model 1 only uses lexical translation
- Translation probability
 - for a foreign sentence $\mathbf{f} = (f_1, \dots, f_{l_f})$ of length l_f
 - to an English sentence $\mathbf{e} = (e_1, \dots, e_{l_e})$ of length l_e
 - with an alignment of each English word e_j to a foreign word f_i according to the alignment function $a : j \rightarrow i$

$$p(\mathbf{e}, a | \mathbf{f}) = \frac{\epsilon}{(l_f + 1)^{l_e}} \prod_{j=1}^{l_e} t(e_j | f_{a(j)})$$

- parameter ϵ is a normalization constant

IBM Model 1 and EM




- EM Algorithm consists of two steps
- Expectation-Step: Apply model to the data
 - parts of the model are hidden (here: alignments)
 - using the model, assign probabilities to possible values
- Maximization-Step: Estimate model from data
 - take assign values as fact
 - collect counts (weighted by probabilities)
 - estimate model from counts
- Iterate these steps until convergence

IBM Model 1 and EM

- Probabilities

$$\begin{aligned} p(\text{the}|\text{la}) &= 0.7 & p(\text{house}|\text{la}) &= 0.05 \\ p(\text{the}|\text{maison}) &= 0.1 & p(\text{house}|\text{maison}) &= 0.8 \end{aligned}$$

- Alignments



$$p(\mathbf{e}, \mathbf{a}|\mathbf{f}) = 0.56 \quad p(\mathbf{e}, \mathbf{a}|\mathbf{f}) = 0.035 \quad p(\mathbf{e}, \mathbf{a}|\mathbf{f}) = 0.08 \quad p(\mathbf{e}, \mathbf{a}|\mathbf{f}) = 0.005$$

$$p(\mathbf{a}|\mathbf{e}, \mathbf{f}) = 0.824 \quad p(\mathbf{a}|\mathbf{e}, \mathbf{f}) = 0.052 \quad p(\mathbf{a}|\mathbf{e}, \mathbf{f}) = 0.118 \quad p(\mathbf{a}|\mathbf{e}, \mathbf{f}) = 0.007$$

- Counts

$$\begin{aligned} c(\text{the}|\text{la}) &= 0.824 + 0.052 & c(\text{house}|\text{la}) &= 0.052 + 0.007 \\ c(\text{the}|\text{maison}) &= 0.118 + 0.007 & c(\text{house}|\text{maison}) &= 0.824 + 0.118 \end{aligned}$$

IBM Model 1 and EM



- We need to be able to compute:
 - Expectation-Step: probability of alignments
 - Maximization-Step: count collection

IBM Model 1 and EM: Expectation Step



- We need to compute $p(a|\mathbf{e}, \mathbf{f})$
- Applying the chain rule:

$$p(a|\mathbf{e}, \mathbf{f}) = \frac{p(\mathbf{e}, a|\mathbf{f})}{p(\mathbf{e}|\mathbf{f})}$$

- We already have the formula for $p(\mathbf{e}, \mathbf{a}|\mathbf{f})$ (definition of Model 1)

IBM Model 1 and EM: Maximization Step



6

- Now we have to collect counts
- Evidence from a sentence pair \mathbf{e}, \mathbf{f} that word e is a translation of word f :

$$c(e|f; \mathbf{e}, \mathbf{f}) = \sum_a p(a|\mathbf{e}, \mathbf{f}) \sum_{j=1}^{l_e} \delta(e, e_j) \delta(f, f_{a(j)})$$

- With the same simplification as before:

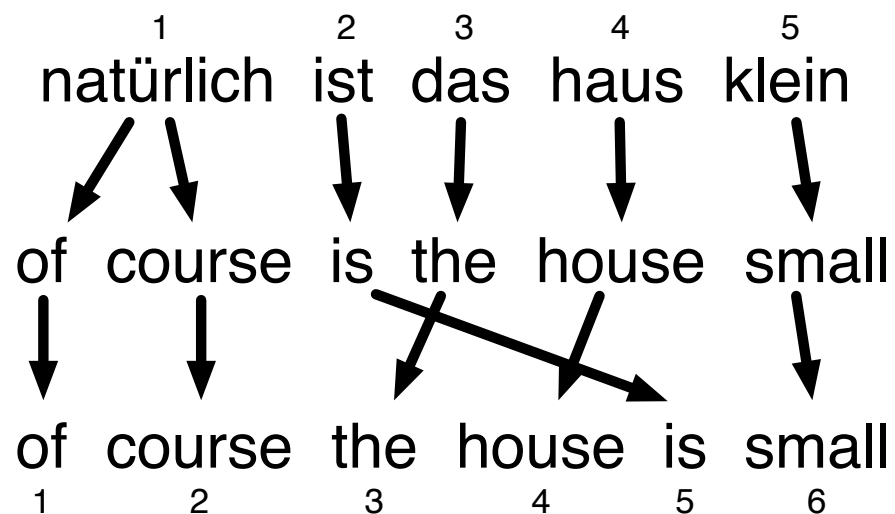
$$c(e|f; \mathbf{e}, \mathbf{f}) = \frac{t(e|f)}{\sum_{i=0}^{l_f} t(e|f_i)} \sum_{j=1}^{l_e} \delta(e, e_j) \sum_{i=0}^{l_f} \delta(f, f_i)$$

ibm model 2

IBM Model 2



Adding a model of alignment



- Modeling alignment with an alignment probability distribution
- Translating English word at position j from foreign word at position $i = a(j)$:

$$a(i|j, l_e, l_f)$$

- Added to IBM Model 1

$$p(\mathbf{e}, a|\mathbf{f}) = \epsilon \prod_{j=1}^{l_e} t(e_j|f_{a(j)}) a(a(j)|j, l_e, l_f)$$

EM Training of IBM Model 2

- Very similar to IBM Model 1 training
 - number of possible word alignments exponential with number of words
 - but: able to reduce complexity of computing $p(\mathbf{e}|\mathbf{f})$ to polynomial
 - same trick applies to IBM Model 2

$$\begin{aligned} p(\mathbf{e}|\mathbf{f}) &= \sum_a p(\mathbf{e}, a|\mathbf{f}) \\ &= \epsilon \sum_{a(1)=0}^{l_f} \dots \sum_{a(l_e)=0}^{l_f} \prod_{j=1}^{l_e} t(e_j|f_{a(j)}) a(a(j)|j, l_e, l_f) \\ &= \epsilon \prod_{j=1}^{l_e} \sum_{i=0}^{l_f} t(e_j|f_{a(j)}) a(a(j)|j, l_e, l_f) \end{aligned}$$

- Count collection for lexical translation

$$c(e|f; \mathbf{e}, \mathbf{f}) = \sum_{j=1}^{l_e} \sum_{i=0}^{l_f} \frac{t(e|f) a(a(j)|j, l_e, l_f) \delta(e, e_j) \delta(f, f_i)}{\sum_{i'=0}^{l_f} t(e|f_{i'}) a(i'|j, l_e, l_f))}$$

- Count collection for alignment

$$c(i|j, l_e, l_f; \mathbf{e}, \mathbf{f}) = \frac{t(e_j|f_i) a(a(j)|j, l_e, l_f)}{\sum_{i'=0}^{l_f} t(e_j|f_{i'}) a(i'|j, l_e, l_f))}$$



- Algorithm for training Model 2 is very similar to the one for IBM Model 1 (pseudo code in book)
- First run a few iterations of IBM Model 1 training
- Initialize probability distributions $t(e|f)$ and $a(i|j, l_e, l_f)$ from IBM Model 1
 - lexical translation probability distribution $t(e|f)$ can be taken verbatim
 - $a(i|j, l_e, l_f)$ initialized to $\frac{1}{l_f+1}$

fast align: reparameterization of ibm model 2

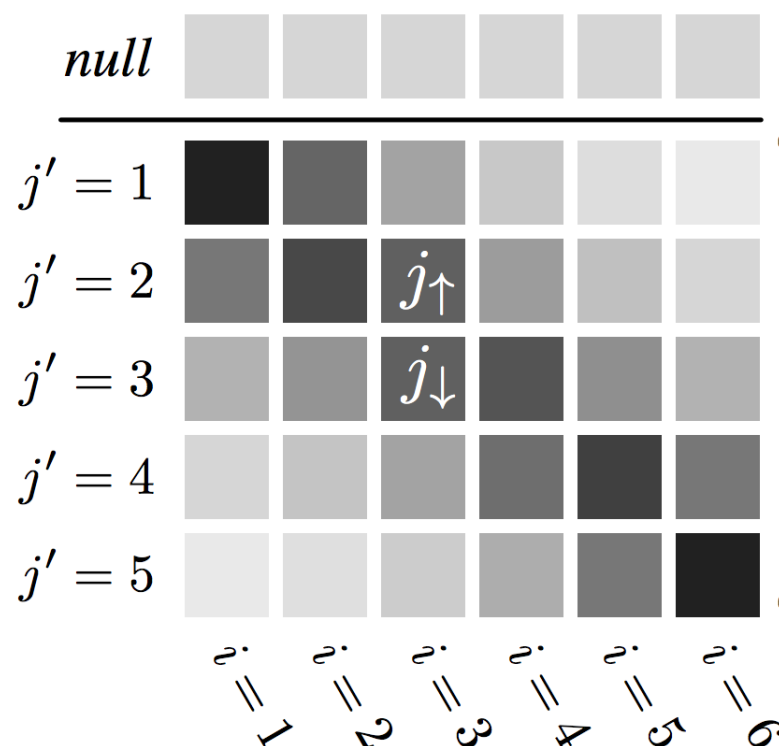
IBM Model 2: A Critique

- Alignment probability distribution has too many parameters ($l_e^2 l_f^2$)

$$a(i|j, l_e, l_f)$$

→ too sparse data to estimate correctly

- Better: bias towards to diagonal



- Distance from diagonal

$$h(i, j, l_e, l_f) = \left| \frac{i}{l_f} - \frac{j}{l_e} \right|$$

- Function that gives higher values to positions close to diagonal (λ is a scaling factor)

$$e^{-\lambda h(i, j, l_e, l_f)}$$

- Special case: alignment to NULL token: p_0
- Alignment probability distribution

$$\delta(a(j) = i | j, l_e, l_f) = \begin{cases} p_0 & \text{if } i = 0 \\ (1 - p_0) \frac{e^{-\lambda h(i, j, l_e, l_f)}}{Z_\lambda(j, m, n)} & \text{if } 0 < i \leq l_e \end{cases}$$

- This model was proposed by Dyer et al. (2013)
- It also changes the word translation probability distribution to include a prior
 - this was originally proposed by Mermer and Saraclar (2011)
 - an efficient estimation method (variational Bayes) was proposed by Riley and Gildea (2012)
- EM training is still simple
 - the probability to align an English word e to a foreign word f does not depend on the choices of other English words
 - the normalization function $Z_\lambda(j, m, n)$ can be computed in $O(1)$

hmm model

- Words do not move independently of each other
 - they often move in groups
 - condition word position on previous word's position
- HMM alignment model:

$$a(a(j)|a(j-1), l_e)$$

- EM algorithm application slightly harder, requires dynamic programming
- IBM Model 4 is similar, also conditions on word classes

EM for the HMM Model

- Main objective: collect fractional counts to estimate
 - word translation probability distribution $t(e_j|f_{a(j)})$
 - alignment probability distribution $a(a(j)|a(j-1), l_e)$
- Consider all possible word alignments
- Collect evidence from each
- Exponentially many \rightarrow need to do this efficiently

Probability of a Word Alignment

	j=1 of	j=2 course	j=3 the	j=4 house	j=5 is	j=6 small
natürlich i=1						
ist i=2						
das i=3						
haus i=4						
klein i=5						

$j = 1$

$$\frac{t(e_1|f_1)}{t(\text{of}|\text{natürlich})}$$

$$\frac{a(a(1)|a(0))}{a(1|0)}$$

$j = 2$

$$\frac{t(e_2|f_1)}{t(\text{course}|\text{natürlich})}$$

$$\frac{a(a(2)|a(1))}{a(1|1)}$$

$j = 3$

$$\frac{t(e_3|f_3)}{t(\text{the}|\text{das})}$$

$$\frac{a(a(3)|a(2))}{a(3|1)}$$

$j = 4$

$$\frac{t(e_4|f_4)}{t(\text{house}|\text{hays})}$$

$$\frac{a(a(4)|a(3))}{a(4|3)}$$

$j = 5$

$$\frac{t(e_5|f_2)}{t(\text{is}|\text{its})}$$

$$\frac{a(a(5)|a(4))}{a(2|4)}$$

$j = 6$

$$\frac{t(e_6|f_5)}{t(\text{small}|\text{klein})}$$

$$\frac{a(a(6)|a(5))}{a(5|2)}$$

	of $j = 1$	course $j = 2$	the $j = 3$...
natürlich $a(j) = 1$	$q_1(1) =$ $t(\text{of} \text{natürlich})$ $\times a(1 0)$			
its $a(j) = 2$	$q_1(2) =$ $t(\text{of} \text{ist})$ $\times a(2 0)$			
das $a(j) = 3$	$q_1(3) =$ $t(\text{of} \text{das})$ $\times a(3 0)$			
...				

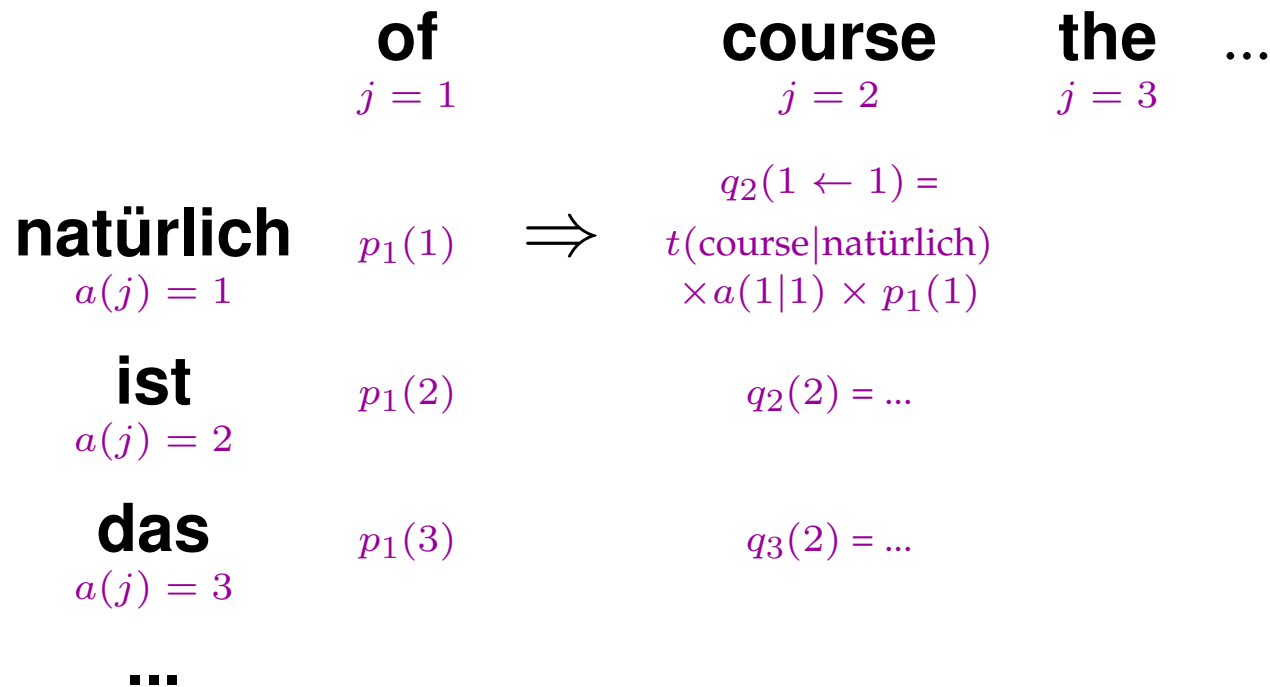
- Compute probabilities for each choice of $i = a(1)$ by normalizing $q_1(i)$

$$p_1(i) = \frac{q_1(i)}{\sum_{i'} q_1(i')}$$

- Use these probabilities for count collection for $t(\text{of}|\bullet)$ and $a(\bullet|0)$

Next English Word

- One way to get there



- Another way to get there

	of $j = 1$	course $j = 2$	the $j = 3$...
natürlich $a(j) = 1$	$p_1(1)$	$q_2(1 \leftarrow 2) =$ $t(\text{course} \text{natürlich})$ $\times a(1 2) \times p_1(2)$		
ist $a(j) = 2$	$p_1(2)$	\Rightarrow	$q_2(2) = \dots$	
das $a(j) = 3$	$p_1(3)$		$q_3(2) = \dots$	
...				

- To compute the score of a state, we have to consider all of the paths

$$q_2(1) = t(e_2|f_1) \times \sum_i p_1(i) a(1|i)$$

Summary of the Math

- Unnormalized score for a transition between two states

$$q_j(i \leftarrow i_{\text{previous}}) = t(e_j|f_i) \times a(i|i_{\text{previous}}) \times p_{j-1}(i_{\text{previous}})$$

- Normalization
$$p_j(i \leftarrow i_{\text{previous}}) = \frac{q_j(i \leftarrow i_{\text{previous}})}{\sum_{i, i_{\text{previous}}} q_j(i \leftarrow i_{\text{previous}})}$$

- Probability of a state
$$p_j(i) = \sum_{i_{\text{previous}}} p_j(i \leftarrow i_{\text{previous}})$$

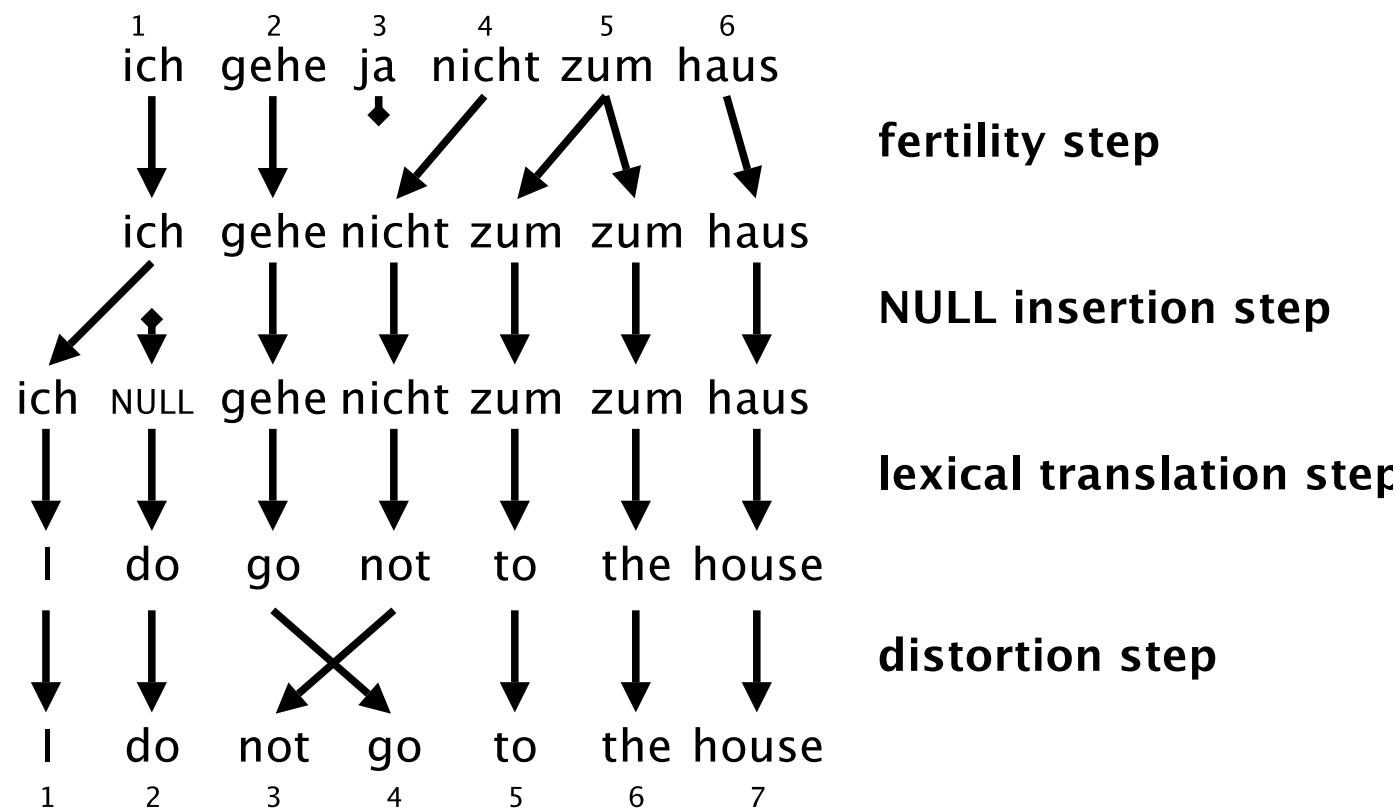
- Count collection
$$c(e_j|f_i) = \sum_{i, j} p_j(i)$$

$$c(i|i_{\text{previous}}) = \sum_{i, j, i_{\text{previous}}} p_j(i \leftarrow i_{\text{previous}})$$

ibm model 3

IBM Model 3

Adding a model of fertility



IBM Model 3: Fertility

- Fertility: number of English words generated by a foreign word
- Modelled by distribution $n(\phi|f)$
- Example:

$$n(1|\text{haus}) \simeq 1$$

$$n(2|\text{zum}) \simeq 1$$

$$n(0|\text{ja}) \simeq 1$$


Sampling the Alignment Space

- Training IBM Model 3 with the EM algorithm
 - The trick that reduces exponential complexity does not work anymore
 - Not possible to exhaustively consider all alignments
- Finding the most probable alignment by hillclimbing
 - start with initial alignment
 - change alignments for individual words
 - keep change if it has higher probability
 - continue until convergence
- Sampling: collecting variations to collect statistics
 - all alignments found during hillclimbing
 - neighboring alignments that differ by a move or a swap

- Better reordering model
- Reordering in IBM Model 2 and 3
 - recall: $d(j|i, l_e, l_f)$
 - for large sentences (large l_f and l_e), sparse and unreliable statistics
 - phrases tend to move together
- Relative reordering model: relative to previously translated words (cepts)

IBM Model 4: Cepts

Foreign words with non-zero fertility forms cepts
(here 5 cepts)



cept π_i	π_1	π_2	π_3	π_4	π_5
foreign position $[i]$	1	2	4	5	6
foreign word $f_{[i]}$	ich	gehe	nicht	zum	haus
English words $\{e_j\}$	I	go	not	to,the	house
English positions $\{j\}$	1	4	3	5,6	7
center of cept \odot_i	1	4	3	6	7

IBM Model 4: Relative Distortion

j	1	2	3	4	5	6	7
e_j	I	do	not	go	to	the	house
in cept $\pi_{i,k}$	$\pi_{1,0}$	$\pi_{0,0}$	$\pi_{3,0}$	$\pi_{2,0}$	$\pi_{4,0}$	$\pi_{4,1}$	$\pi_{5,0}$
\odot_{i-1}	0	-	4	1	3	-	6
$j - \odot_{i-1}$	+1	-	-1	+3	+2	-	+1
distortion	$d_1(+1)$	1	$d_1(-1)$	$d_1(+3)$	$d_1(+2)$	$d_{>1}(+1)$	$d_1(+1)$

- Center \odot_i of a cept π_i is $\text{ceiling}(\text{avg}(j))$
- Three cases:
 - uniform for NULL generated words
 - first word of a cept: d_1
 - next words of a cept: $d_{>1}$

- Some words may trigger reordering → condition reordering on words

for initial word in cept: $d_1(j - \odot_{[i-1]} | f_{[i-1]}, e_j)$

for additional words: $d_{>1}(j - \Pi_{i,k-1} | e_j)$

- Sparse data concerns → cluster words into classes

for initial word in cept: $d_1(j - \odot_{[i-1]} | \mathcal{A}(f_{[i-1]}), \mathcal{B}(e_j))$

for additional words: $d_{>1}(j - \Pi_{i,k-1} | \mathcal{B}(e_j))$

- IBM Models 1–4 are *deficient*
 - some impossible translations have positive probability
 - multiple output words may be placed in the same position
 - probability mass is wasted
- IBM Model 5 fixes deficiency by keeping track of vacancies (available positions)

- IBM Models were the pioneering models in statistical machine translation
- Introduced important concepts
 - generative model
 - EM training
 - reordering models
- Only used for niche applications as translation model
- ... but still in common use for word alignment (e.g., GIZA++ toolkit)

word alignment

Word Alignment

Given a sentence pair, which words correspond to each other?







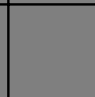

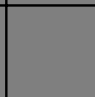

	michael	geht	davon	aus	,	dass	er	im	haus	bleibt
michael										
assumes										
that										
he										
will										
stay										
in										
the										
house										

Word Alignment?

	john	wohnt	hier	nicht
john				
does		?		?
not				
live				
here				

Is the English word **does** aligned to the German **wohnt** (verb) or **nicht** (negation) or neither?

Word Alignment?

	john	biss	ins	grass
john				
kicked				
the				
bucket				

How do the idioms **kicked the bucket** and **biss ins grass** match up?
Outside this exceptional context, **bucket** is never a good translation for **grass**

Measuring Word Alignment Quality

- Manually align corpus with *sure* (S) and *possible* (P) alignment points ($S \subseteq P$)
- Common metric for evaluation word alignments: Alignment Error Rate (AER)

$$\text{AER}(S, P; A) = \frac{|A \cap S| + |A \cap P|}{|A| + |S|}$$

- $\text{AER} = 0$: alignment A matches all sure, any possible alignment points
- However: different applications require different precision/recall trade-offs

symmetrization

Word Alignment with IBM Models

- IBM Models create a **many-to-one** mapping
 - words are aligned using an alignment function
 - a function may return the same value for different input (one-to-many mapping)
 - a function can not return multiple values for one input (no many-to-one mapping)
- Real word alignments have **many-to-many** mappings

- Run IBM Model training in both directions

→ two sets of word alignment points

- Intersection: high precision alignment points
- Union: high recall alignment points
- Refinement methods explore the sets between intersection and union

Example

english to spanish

	bofetada				bruja			
	Maria	no	daba	una	a	la	verde	
Mary	■							
did					■			
not		■						
slap			■	■	■			
the						■		
green								■
witch							■	


spanish to english

	bofetada				bruja			
	Maria	no	daba	una	a	la	verde	
Mary	■							
did		■						
not		■						
slap					■			
the						■		
green								■
witch							■	

intersection

	bofetada				bruja			
	Maria	no	daba	una	a	la	verde	
Mary	■							
did								
not		■						
slap					■			
the						■		
green								■
witch							■	

Growing Heuristics



black: intersection

grey: additional points in union

- Add alignment points from union based on heuristics:
 - directly/diagonally neighboring points
 - finally, add alignments that connect unaligned words in source and/or target
- Popular method: grow-diag-final-and

grow-diag-final(e2f,f2e)

- 1: neighboring = $\{(-1,0),(0,-1),(1,0),(0,1),(-1,-1),(-1,1),(1,-1),(1,1)\}$
- 2: alignment $A = \text{intersect}(e2f,f2e)$; **grow-diag()**; **final(e2f)**; **final(f2e)**;

grow-diag()

- 1: **while** new points added **do**
- 2: **for all** English word $e \in [1 \dots e_n]$, foreign word $f \in [1 \dots f_n]$, $(e, f) \in A$ **do**
- 3: **for all** neighboring alignment points $(e_{\text{new}}, f_{\text{new}})$ **do**
- 4: **if** $(e_{\text{new}}$ unaligned OR f_{new} unaligned) AND $(e_{\text{new}}, f_{\text{new}}) \in \text{union}(e2f,f2e)$ **then**
- 5: add $(e_{\text{new}}, f_{\text{new}})$ to A
- 6: **end if**
- 7: **end for**
- 8: **end for**
- 9: **end while**

final()

- 1: **for all** English word $e_{\text{new}} \in [1 \dots e_n]$, foreign word $f_{\text{new}} \in [1 \dots f_n]$ **do**
- 2: **if** $(e_{\text{new}}$ unaligned OR f_{new} unaligned) AND $(e_{\text{new}}, f_{\text{new}}) \in \text{union}(e2f,f2e)$ **then**
- 3: add $(e_{\text{new}}, f_{\text{new}})$ to A
- 4: **end if**
- 5: **end for**

More Work on Symmetrization

- Symmetrize after each iteration of IBM Models [Matusov et al., 2004]
 - run one iteration of E-step for each direction
 - symmetrize the two directions
 - count collection (M-step)
- Use of posterior probabilities in symmetrization
 - generate n-best alignments for each direction
 - calculate how often an alignment point occurs in these alignments
 - use this posterior probability during symmetrization

- Link deletion [Fossum et al., 2008]
 - start with union of IBM Model alignment points
 - delete one alignment point at a time
 - uses a neural network classifiers that also considers aspects such as how useful the alignment is for learning translation rules
- Link addition [Ren et al., 2007] [Ma et al., 2008]
 - possibly start with a skeleton of highly likely alignment points
 - add one alignment point at a time

- Given some annotated training data, supervised learning methods are possible
- Structured prediction
 - not just a classification problem
 - solution structure has to be constructed in steps
- Many approaches: maximum entropy, neural networks, support vector machines, conditional random fields, MIRA, ...
- Small labeled corpus may be used for parameter tuning of unsupervised aligner [Fraser and Marcu, 2007]

- Aligning phrases
 - joint model [Marcu and Wong, 2002]
 - problem: EM algorithm likes really long phrases
- Fraser's LEAF
 - decomposes word alignment into many steps
 - similar in spirit to IBM Models
 - includes step for grouping into phrase
- Riesa's NILE
 - use syntactic parse trees to guide word alignment
 - build up words bottom up following the parse tree

Final Remarks

- Research on word alignment has recently picked up again
 - speed matters
 - incremental (“online”) training
- Unclear link between
 - word alignment quality measured against manual gold standard
 - impact on machine translation quality
- Advice: if you develop method, make easy-to-use toolkit available